

shaikat_303527_exercise_6-Copy1

December 4, 2019

1 Regression Datasets

1.0.1 Generating D1 dataset

```
In [2]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import SGDRegressor
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import GridSearchCV
import math as Math
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.preprocessing import PolynomialFeatures
from sklearn.exceptions import DataConversionWarning
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=DataConversionWarning)
warnings.filterwarnings("ignore", category=DeprecationWarning)

d1_x=np.random.uniform(1,0.005,size=(100,1))
r=np.random.rand(100,1)
d1_y=1.3*(d1_x**2)+4.8*d1_x+8+r
d1_x = (d1_x - d1_x.mean())/d1_x.std() #data normalized
# x_train_d1, x_test_d1, y_train_d1, y_test_d1 =train_test_split(x,d1_y,train_size=0.8
# test_size=0.2
# random_state=
```

1.0.2 Preprocessing Wine Quality Red dataset

```
In [3]: filename=r"E:\Documents\University of Hildesheim\Machine learning lab\lab5\winequality
        rwine_data = pd.read_csv(filename,delimiter=';')
        rwine_data.head(3)
```

```
Out[3]:   fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \
0           7.4                0.70           0.00              1.9       0.076
```

1	7.8	0.88	0.00	2.6	0.098
2	7.8	0.76	0.04	2.3	0.092

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	11.0	34.0	0.9978	3.51	0.56
1	25.0	67.0	0.9968	3.20	0.68
2	15.0	54.0	0.9970	3.26	0.65

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5

The data has no numeric values

```
In [4]: rwine_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
fixed acidity          1599 non-null float64
volatile acidity       1599 non-null float64
citric acid            1599 non-null float64
residual sugar         1599 non-null float64
chlorides              1599 non-null float64
free sulfur dioxide    1599 non-null float64
total sulfur dioxide   1599 non-null float64
density                1599 non-null float64
pH                    1599 non-null float64
sulphates              1599 non-null float64
alcohol                1599 non-null float64
quality                1599 non-null int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
In [5]: Xdata_rwine = rwine_data.loc[:,rwine_data.columns!='quality']
        Ydata_rwine = rwine_data[['quality']]
        Xdata_rwine = (Xdata_rwine - Xdata_rwine.mean())/Xdata_rwine.std() #data normalized
        x_train_rwine, x_test_rwine, y_train_rwine, y_test_rwine =train_test_split(Xdata_rwine,
                                                                                      Ydata_rwine,tra
                                                                                      test_size=0.2,
                                                                                      random_state=0)
```

```
In [6]: y_train_rwine=y_train_rwine.values.reshape(-1,1)
        y_test_rwine=y_test_rwine.values.reshape(-1,1)
        x_train_rwine=x_train_rwine.values
        x_test_rwine=x_test_rwine.values
```

2 GLMs


2.0.1 Generalized Linear Models with Scikit Learn

2.0.2 1. Ordinary least Squares

```
In [7]: learning_rate=[1,0.8,0.5,0.1,0.001,0.00001]
        rmsetrain=[]
        rmsetest=[]
        for i in learning_rate:
            model=SGDRegressor(eta0=i,learning_rate='constant',penalty=None,shuffle=True)
            model.fit(x_train_rwine,y_train_rwine)
            y_pred=model.predict(x_train_rwine)
            y_pred_test=model.predict(x_test_rwine)
            rmsetrain.append(Math.sqrt(mean_squared_error(y_train_rwine,y_pred)))
            rmsetest.append(Math.sqrt(mean_squared_error(y_test_rwine,y_pred_test)))
```

2.0.3 Analyzing the graph we can see that the error increases as the learning rate increases


```
In [8]: plt.plot(learning_rate,rmsetrain,'g',label='train_sample')
        plt.plot(learning_rate,rmsetest,'r',label='test_sample')
        plt.xlabel('learning rate')
        plt.ylabel('RMSE')
        plt.legend()
        plt.show()
        plt.close()
```




2.0.4 2. Ridge Regression

```
In [8]: hyp_par=([0.8,1],[0.5,0.1],[0.001,0.01])
        rmsetrain=[]
        rmsetest=[]
        alpha_com=[]
        lamda_com=[]
        for i,j in hyp_par:
            model=SGDRegressor(eta0=i,learning_rate='constant',penalty='l2',shuffle=True,alpha=0.001)
            model.fit(x_train_rwine,y_train_rwine)
            y_pred=model.predict(x_train_rwine)
            y_pred_test=model.predict(x_test_rwine)
            rmsetrain.append(Math.sqrt(mean_squared_error(y_train_rwine,y_pred)))
            rmsetest.append(Math.sqrt(mean_squared_error(y_test_rwine,y_pred_test)))
            alpha_com.append(i)
            lamda_com.append(j)

In [9]: from mpl_toolkits import mplot3d
        fig = plt.figure()
        fig.set_figheight(10)
        fig.set_figwidth(10)
        ax = plt.axes(projection= '3d')
        ax.scatter3D(alpha_com, lamda_com, rmsetrain, c=rmsetrain, cmap='hsv')
        ax.set_xlabel('alpha')
        ax.set_ylabel('lamda')
        ax.set_zlabel('RMSE_TRAIN')
        plt.show()
```



```
In [10]: fig = plt.figure()
fig.set_figheight(10)
fig.set_figwidth(10)
ax = plt.axes(projection= '3d')
ax.scatter3D(alpha_com, lamda_com, rmsetest, c=rmsetest, cmap='hsv')
ax.set_xlabel('alpha')
ax.set_ylabel('lamda')
ax.set_zlabel('RMSE_TEST')
plt.show()
```



2.0.5 2. LASSO

```
In [11]: hyp_par=([0.8,1],[0.5,0.1],[0.001,0.01])
          rmsetrain=[]
          rmsetest=[]
          alpha_com=[]
          lamda_com=[]
          for i,j in hyp_par:
              model=SGDRegressor(eta0=i,learning_rate='constant',penalty='l1',shuffle=True,alpha
              model.fit(x_train_rwine,y_train_rwine)
              y_pred=model.predict(x_train_rwine)
              y_pred_test=model.predict(x_test_rwine)
              rmsetrain.append(Math.sqrt(mean_squared_error(y_train_rwine,y_pred)))
              rmsetest.append(Math.sqrt(mean_squared_error(y_test_rwine,y_pred_test)))
```

```


alpha_com.append(i)
lamda_com.append(j)

```

```

In [12]: fig = plt.figure()
fig.set_figheight(10)
fig.set_figwidth(10)
ax = plt.axes(projection= '3d')
ax.scatter3D(alpha_com, lamda_com, rmsetrain, c=rmsetrain, cmap='hsv')
ax.set_xlabel('alpha')
ax.set_ylabel('lamda')
ax.set_zlabel('RMSE_TRAIN')
plt.show()

```



```


In [13]: fig = plt.figure()
fig.set_figheight(10)

```

```

fig.set_figwidth(10)
ax = plt.axes(projection= '3d')
ax.scatter3D(alpha_com, lamda_com, rmsetrain, c=rmsetest, cmap='hsv')
ax.set_xlabel('alpha')
ax.set_ylabel('lamda')
ax.set_zlabel('RMSE_TEST')
plt.show()

```



2.0.6 A standard least squares model won't generalize well for a data set different than its training data. By increasing the bias regularization reduces the variance in the dataset which shown in the 2d and 3d plots. Hence, the graphs show that regularization reduces overfitting of the model.


3 Hyperparameters using scikit learn GridSearchCV

3.0.1 Hyperparameters tuning of Ordinary Least Squares

```
In [14]: hyp_par={'eta0':[1,0.8,0.5,0.1,0.001,0.00001]}
         model=SGDRegressor(learning_rate='constant',shuffle=True,penalty=None)
         gs_ols=GridSearchCV(model,hyp_par,cv=5)
         gs_ols.fit(x_train_rwine,y_train_rwine)
```

```
Out[14]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=SGDRegressor(alpha=0.0001, average=False, early_stopping=False, epsi
                      eta0=0.01, fit_intercept=True, l1_ratio=0.15,
                      learning_rate='constant', loss='squared_loss', max_iter=None,
                      n_iter=None, n_iter_no_change=5, penalty=None, power_t=0.25,
                      random_state=None, shuffle=True, tol=None, validation_fraction=0.1,
                      verbose=0, warm_start=False),
                      fit_params=None, iid='warn', n_jobs=None,
                      param_grid={'eta0': [1, 0.8, 0.5, 0.1, 0.001, 1e-05]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring=None, verbose=0)
```

```
In [15]: plt.plot(hyp_par['eta0'],gs_ols.cv_results_['mean_train_score'],'g',label='train_samp
         plt.plot(hyp_par['eta0'],gs_ols.cv_results_['mean_test_score'],'r',label='test_sample
         plt.xlabel('learning rate')
         plt.ylabel('mean score')
         plt.legend()
         plt.show()
         plt.close()
```



3.0.2 Hyperparameters tuning of Ridge Regression

```
In [16]: hyp_par={'eta0':[1,0.8,0.5,0.1,0.001,0.00001], 'alpha':[0.8,0.5,0.3,0.1,0.001,0.0001]}
```

```
model=SGDRegressor(learning_rate='constant',shuffle=True,penalty='l1')
gs_rr=GridSearchCV(model,hyp_par,cv=5)
gs_rr.fit(x_train_rwine,y_train_rwine)
```


```
Out[16]: GridSearchCV(cv=5, error_score='raise-deprecating',
    estimator=SGDRegressor(alpha=0.0001, average=False, early_stopping=False, epsi
    eta0=0.01, fit_intercept=True, l1_ratio=0.15,
    learning_rate='constant', loss='squared_loss', max_iter=None,
    n_iter=None, n_iter_no_change=5, penalty='l1', power_t=0.25,
    random_state=None, shuffle=True, tol=None, validation_fraction=0.1,
    verbose=0, warm_start=False),
    fit_params=None, iid='warn', n_jobs=None,
    param_grid={'eta0': [1, 0.8, 0.5, 0.1, 0.001, 1e-05], 'alpha': [0.8, 0.5, 0.3,
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring=None, verbose=0)
```

```
In [17]: alpha_com=[]
    lamda_com=[]
    for i in range(len(gs_rr.cv_results_['params'])):
        alpha_com.append(gs_rr.cv_results_['params'][i]['alpha'])
        lamda_com.append(gs_rr.cv_results_['params'][i]['eta0'])
```

```

In [18]: fig = plt.figure()
fig.set_figheight(10)
fig.set_figwidth(10)
ax = plt.axes(projection= '3d')
ax.scatter3D(alpha_com, lamda_com, gs_rr.cv_results_['mean_test_score'], c=gs_rr.cv_r
ax.set_xlabel('alpha')
ax.set_ylabel('lamda')
ax.set_zlabel('mean_test_score')
plt.show()

```



```


In [19]: fig = plt.figure()
fig.set_figheight(10)
fig.set_figwidth(10)
ax = plt.axes(projection= '3d')

```

```

ax.scatter3D(alpha_com, lamda_com, gs_rr.cv_results_['mean_train_score'], c=gs_rr.cv_
ax.set_xlabel('alpha')
ax.set_ylabel('lamda')
ax.set_zlabel('mean_train_score')
plt.show()

```



3.0.3 Hyperparameters tuning of LASSO

```

In [20]: hyp_par={'eta0':[1,0.8,0.5,0.1,0.001,0.00001], 'alpha':[0.8,0.5,0.3,0.1,0.001,0.0001]}

model=SGDRegressor(learning_rate='constant',shuffle=True,penalty='l2')
gs_lasso=GridSearchCV(model,hyp_par,cv=5)
gs_lasso.fit(x_train_rwine,y_train_rwine)

```

```

Out[20]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=SGDRegressor(alpha=0.0001, average=False, early_stopping=False, epsi
                      eta0=0.01, fit_intercept=True, l1_ratio=0.15,
                      learning_rate='constant', loss='squared_loss', max_iter=None,
                      n_iter=None, n_iter_no_change=5, penalty='l2', power_t=0.25,
                      random_state=None, shuffle=True, tol=None, validation_fraction=0.1,
                      verbose=0, warm_start=False),
                      fit_params=None, iid='warn', n_jobs=None,
                      param_grid={'eta0': [1, 0.8, 0.5, 0.1, 0.001, 1e-05], 'alpha': [0.8, 0.5, 0.3,
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring=None, verbose=0)

```

```


In [21]: alpha_com=[]
         lamda_com=[]
         for i in range(len(gs_lasso.cv_results_['params'])):
             alpha_com.append(gs_lasso.cv_results_['params'][i]['alpha'])
             lamda_com.append(gs_lasso.cv_results_['params'][i]['eta0'])

```


```

In [22]: fig = plt.figure()
         fig.set_figheight(10)
         fig.set_figwidth(10)
         ax = plt.axes(projection= '3d')
         ax.scatter3D(alpha_com, lamda_com, gs_lasso.cv_results_['mean_test_score'], c=gs_lasso
         ax.set_xlabel('alpha')
         ax.set_ylabel('lamda')
         ax.set_zlabel('mean_test_score')
         plt.show()

```



```
In [23]: fig = plt.figure()
fig.set_figheight(10)
fig.set_figwidth(10)
ax = plt.axes(projection= '3d')
ax.scatter3D(alpha_com, lamda_com, gs_lasso.cv_results_['mean_train_score'], c=gs_lasso.cv_results_['mean_test_score'])
ax.set_xlabel('alpha')
ax.set_ylabel('lamda')
ax.set_zlabel('mean_train_score')
plt.show()
```



3.0.4 The tuning parameter 'alpha' controls the impact on bias and variance. As the value of alpha rises, it reduces the value of coefficients and thus reducing the variance. But after certain value of alpha, there is a rise to bias in the model and thus underfitting. Therefore, the value of alpha should be carefully selected.

3.0.5 Cross validation on ordinary least squares

```
In [24]: model=SGDRegressor(learning_rate='constant',shuffle=True,penalty=None,eta0=gs_ols.best_
model.fit(x_train_rwine,y_train_rwine)
y_pred=model.predict(x_train_rwine)
y_pred_test=model.predict(x_test_rwine)
cross_val_train_ols=cross_val_score(model,x_train_rwine,y_train_rwine,cv=5)
cross_val_test_ols=cross_val_score(model,x_test_rwine,y_test_rwine,cv=5)
```

3.0.6 Cross validation on Ridge Regression

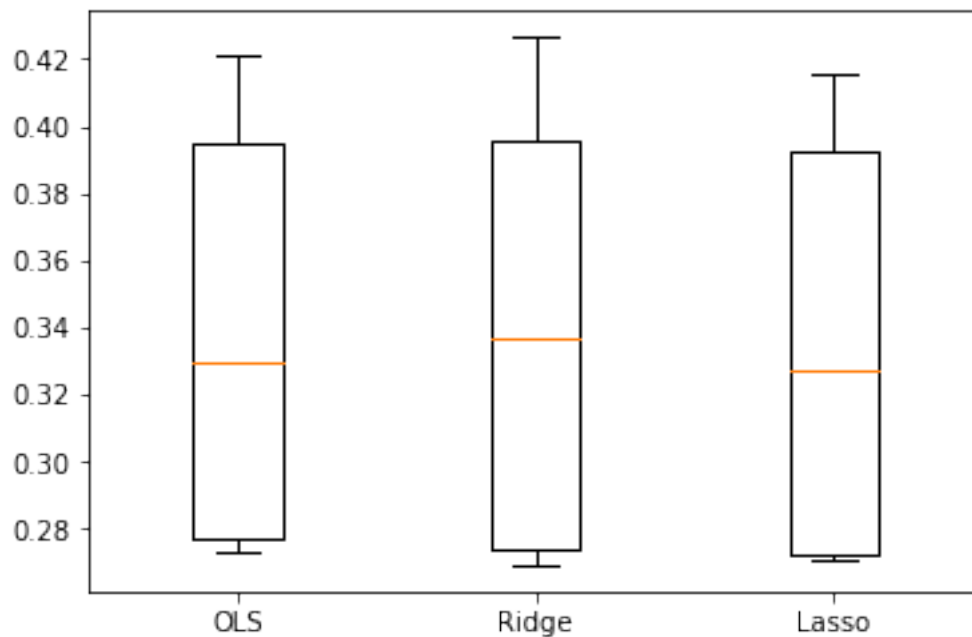
```
In [25]: model=SGDRegressor(learning_rate='constant',shuffle=True,penalty='l2',eta0=gs_rr.best.  
model.fit(x_train_rwine,y_train_rwine)  
y_pred=model.predict(x_train_rwine)  
y_pred_test=model.predict(x_test_rwine)  
cross_val_train_rr=cross_val_score(model,x_train_rwine,y_train_rwine,cv=5)  
cross_val_test_rr=cross_val_score(model,x_test_rwine,y_test_rwine,cv=5)
```

3.0.7 Cross validation on LASSO

```
In [26]: model=SGDRegressor(learning_rate='constant',shuffle=True,penalty='l2',eta0=gs_lasso.b  
model.fit(x_train_rwine,y_train_rwine)  
y_pred=model.predict(x_train_rwine)  
y_pred_test=model.predict(x_test_rwine)  
cross_val_train_lasso=cross_val_score(model,x_train_rwine,y_train_rwine,cv=5)  
cross_val_test_lasso=cross_val_score(model,x_test_rwine,y_test_rwine,cv=5)
```

```
In [27]: print('Demonstrating training data of all the models in box plot')  
  
plt.boxplot([cross_val_train_ols,cross_val_train_rr,cross_val_train_lasso])  
plt.xticks(np.arange(1,4),('OLS','Ridge','Lasso'))  
plt.show()
```


Demonstrating training data of all the models in box plot




```
In [28]: print('Demonstrating test data of all the models in box plot')

plt.boxplot([cross_val_test_ols, cross_val_test_rr, cross_val_test_lasso])
plt.xticks(np.arange(1,4), ('OLS', 'Ridge', 'Lasso'))
plt.show()
```

Demonstrating test data of all the models in box plot



3.0.8 After doing the cross validation score and representing it in a box plot we can see that regularized model is better than unregularized model

4 Polynomial Regression


4.0.1 Prediction with high degree of polynomials

```
In [29]: fig, axes = plt.subplots(1, 1, figsize=(10, 10))
axes.scatter(d1_x, d1_y, color='orange', marker='*', label="training points")
axes.set_xlabel('X')
axes.set_ylabel('predicted Y')
axes.legend()
rmse = []
degree = [1, 2, 7, 10, 16, 100]
for i in degree:
    model = LinearRegression()
    polyf = PolynomialFeatures(degree=i)
```

```

poly_x=polyf.fit_transform(d1_x)
model.fit(poly_x,d1_y)
y_pred=model.predict(poly_x)
X,Y=zip(*sorted(zip(d1_x,y_pred)))
axs.plot(X,Y,label="degree %i" % i)
axs.legend()
plt.show()

```



4.0.2 As the degree increases the error also increases which is illustrated in the graph

4.0.3 Effect of Regularization

```


In [30]: lamda=[0,10**-6,10**-2,1]
fig,axs=plt.subplots(1,1,figsize=(10,10))

```

```

axs.scatter(d1_x,d1_y,color='orange',marker='*',label="training points")
axs.set_xlabel('X')
axs.set_ylabel('predicted Y')
axs.legend()
for i in lamda:
    model=Ridge(alpha=i)
    polyf=PolynomialFeatures(degree=10)
    poly_x=polyf.fit_transform(d1_x)
    model.fit(poly_x,d1_y)
    y_pred=model.predict(poly_x)
    X,Y=zip(*sorted(zip(d1_x,y_pred)))
    axs.plot(X,Y)
axs.legend()
plt.show()

```



4.0.4 We can see that lower lamda values are giving good results with the degree=10

5 Coordinate Descent

5.0.1 Implementing Coordinate Descent

```
In [202]: def coordinate_descent(x,y,epochs):
            m_train,n_features = np.shape(x)
            beta
                = np.zeros(n_features).reshape(-1,1)

            beta_hist=np.zeros((epochs+1,n_features))
            hist = []
            for j in range(epochs):
                for i in range(len(beta)):
                    # choosing x coordinate to update
                    _x = x[:,i]
                    x_coor=np.delete(x,i,axis=1)

                    #choosing beta to update
                    beta_coor=np.delete(beta,i,axis=0)


                    col = x_coor.dot(beta_coor)
                    num = ((y - col).T).dot(_x)
                    den = _x.T.dot(_x)
                    update = num/den
                    beta[i] = update
                    #print(update)

                beta_hist[j+1] = beta.ravel()

            return beta,beta_hist

In [203]: mtrain,n = x_train_rwine.shape
            Xtrain = np.concatenate((x_train_rwine, np.ones(mtrain).reshape(-1,1)), axis = 1)

In [204]: epochs=30
            plt.figure(figsize = (15,12))
            beta,hist = coordinate_descent(Xtrain,y_train_rwine,epochs)
            plt.plot(hist)
            plt.show()
            plt.close()
```



```
In [205]: def soft_threshold(beta,reg):
    if beta < - reg:
        return (beta + reg)
    elif beta > reg:
        return (beta - reg)
    else:
        return 0

def coordinate_descent(x,y,epochs,lamda = 1, lasso = False):
    m_train,n_features = np.shape(x)
    beta = np.zeros(n_features).reshape(-1,1)

    beta_hist=np.zeros((epochs+1,n_features))
    hist = []

    if lasso:
        print(f"Using Lasso with regularization {lamda}")

    for j in range(epochs):
        for i in range(len(beta)):
```

```

    # choosing x coordinate to update
    _x = x[:,i]
    x_coor=np.delete(x,i,axis=1)

    #choosing beta to update
    beta_coor=np.delete(beta,i,axis=0)

    col = x_coor.dot(beta_coor)
    num = ((y - col).T).dot(_x)
    den = _x.T.dot(_x)
    update = num/den

    if lasso:
        reg = 0.5 * lamda / (_x.T.dot(_x))
        update = soft_threshold(update,reg)

    beta[i] = update
    #print(update)

    beta_hist[j+1] = beta.ravel()

    return beta,beta_hist

```

```

In [206]: mtrain,n = x_train_rwine.shape
          Xtrain = np.concatenate((x_train_rwine, np.ones(mtrain).reshape(-1,1)), axis = 1)


```

```

In [208]: epochs=30
          plt.figure(figsize = (15,12))
          beta,hist = coordinate_descent(Xtrain,y_train_rwine,epochs,lamda = 0.000001, lasso =
          plt.plot(hist)
          plt.show()
          plt.close()

```

Using Lasso with regularization 1e-06



5.0.2 comparing the two graphs we can see that the regularized model betas are more closer to each other than regularized model

In []: